

```
"""
```

```
1. Implement and demonstrate the FIND-S algorithm for finding the
```

```
most specific hypothesis based on a given set of training data samples.
```

```
Read the training data from a .CSV file
```

```
"""
```

```
import csv
```

```
a = []
```

```
with open('enjoysport.csv', 'r') as csvfile:
```

```
    next(csvfile)
```

```
    for row in csv.reader(csvfile):
```

```
        a.append(row)
```

```
    print(a)
```

```
print("\n\nThe total number of training instances are : ",len(a))
```

```
num_attribute = len(a[0])-1
```

```
print("\n\nThe initial hypothesis is : ")
```

```
hypothesis = ['0']*num_attribute
```

```
print(hypothesis)
```

```
for i in range(0, len(a)):
```

```
    if a[i][num_attribute] == 'yes':
```

```

print ("\nInstance ", i+1, "is", a[i], " and is Positive Instance")

for j in range(0, num_attribute):

    if hypothesis[j] == '0' or hypothesis[j] == a[i][j]:

        hypothesis[j] = a[i][j]

    else:

        hypothesis[j] = '?'

print("The hypothesis for the training instance", i+1, " is: " , hypothesis, "\n")

if a[i][num_attribute] == 'no':

    print ("\nInstance ", i+1, "is", a[i], " and is Negative Instance Hence Ignored")

    print("The hypothesis for the training instance", i+1, " is: " , hypothesis, "\n")

print("\nThe Maximally specific hypothesis for the training instance is ", hypothesis)

```

Output:



```

Python 3.10.4 (tags/v3.10.4:9d38120, Mar 23 2022, 23:13:41) [MSC v.1929 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
===== RESTART: C:\machine learning programs\lab\sample.py =====
[['sunny', 'warm', 'normal', 'strong', 'warm', 'same', 'yes'], ['sunny', 'warm', 'high', 'strong', 'warm', 'same', 'yes'], ['rainy', 'cold', 'high', 'strong', 'warm', 'change', 'no'], ['sunny', 'warm', 'high', 'strong', 'cool', 'change', 'yes']]

The total number of training instances are : 4

The initial hypothesis is :
['0', '0', '0', '0', '0', '0']

Instance 1 is ['sunny', 'warm', 'normal', 'strong', 'warm', 'same', 'yes'] and is Positive Instance
The hypothesis for the training instance 1 is: ['sunny', 'warm', 'normal', 'strong', 'warm', 'same']

Instance 2 is ['sunny', 'warm', 'high', 'strong', 'warm', 'same', 'yes'] and is Positive Instance
The hypothesis for the training instance 2 is: ['sunny', 'warm', '?', 'strong', 'warm', 'same']

Instance 3 is ['rainy', 'cold', 'high', 'strong', 'warm', 'change', 'no'] and is Negative Instance Hence Ignored
The hypothesis for the training instance 3 is: ['sunny', 'warm', '?', 'strong', 'warm', 'same']

Instance 4 is ['sunny', 'warm', 'high', 'strong', 'cool', 'change', 'yes'] and is Positive Instance
The hypothesis for the training instance 4 is: ['sunny', 'warm', '?', 'strong', '?', '?']

The Maximally specific hypothesis for the training instance is ['sunny', 'warm', '?', 'strong', '?', '?']
>>>

```

- For a given set of training data examples stored in a .CSV file, implement and demonstrate the Candidate-Elimination algorithm in python to output a description of the set of all hypotheses consistent with the training examples.

```
import numpy as np
import pandas as pd

data = pd.read_csv(path+'/enjoysport.csv')
concepts = np.array(data.iloc[:,0:-1])
print("\nInstances are:\n",concepts)
target = np.array(data.iloc[:,-1])
print("\nTarget Values are: ",target)

def learn(concepts, target):
    specific_h = concepts[0].copy()
    print("\nInitialization of specific_h and general_h")
    print("\nSpecific Boundary: ", specific_h)
    general_h = [["?" for i in range(len(specific_h))] for i in range(len(specific_h))]
    print("\nGeneric Boundary: ",general_h)

    for i, h in enumerate(concepts):
        print("\nInstance", i+1, "is ", h)
        if target[i] == "yes":
            print("Instance is Positive ")
            for x in range(len(specific_h)):
                if h[x] != specific_h[x]:
                    specific_h[x] = '?'
                    general_h[x][x] = '?'

        if target[i] == "no":
            print("Instance is Negative ")
            for x in range(len(specific_h)):
                if h[x] != specific_h[x]:
                    general_h[x][x] = specific_h[x]
                else:
                    general_h[x][x] = '?'

    print("Specific Boundary after ", i+1, "Instance is ", specific_h)
    print("Generic Boundary after ", i+1, "Instance is ", general_h)
    print("\n")
```

Output:

```
===== RESTART: C:

Instances are:
[['sunny' 'warm' 'normal' 'strong' 'warm' 'same']
 ['sunny' 'warm' 'high' 'strong' 'warm' 'same']
 ['rainy' 'cold' 'high' 'strong' 'warm' 'change']
 ['sunny' 'warm' 'high' 'strong' 'cool' 'change']]

Target Values are:  ['yes' 'yes' 'no' 'yes']
> |
```

- Write a program to demonstrate the working of the decision tree based ID3 algorithm. Use an appropriate data set for building the decision tree and apply this knowledge to classify a new sample.

#Import libraries and read data using read_csv() function. Remove the target from the data and store #attributes in the features variable.

```
import pandas as pd
import math
import numpy as np
```

```
data = pd.read_csv("Dataset/4-dataset.csv")
```

```
features = [feat for feat in data]
```

```
features.remove("answer")
```

Create a class named Node with four members children, value, isLeaf and pred.

```
class Node:
```

```
    def __init__(self):
```

```
        self.children = []
```

```
        self.value = ""
```

```
        self.isLeaf = False
```

```
        self.pred = ""
```

Define a function called entropy to find the entropy of the dataset

```
def entropy(examples):
```

```
    pos = 0.0
```

```
    neg = 0.0
```

```
    for _, row in examples.iterrows():
```

```
        if row["answer"] == "yes":
```

```
            pos += 1
```

```
        else:
```

```
            neg += 1
```

```
    if pos == 0.0 or neg == 0.0:
```

```
        return 0.0
```

```
    else:
```

```
        p = pos / (pos + neg)
```

```
        n = neg / (pos + neg)
```

```
        return -(p * math.log(p, 2) + n * math.log(n, 2))
```

Define a function named info_gain to find the gain of the attribute

```
def info_gain(examples, attr):
```

```
    uniq = np.unique(examples[attr])
```

```
    #print ("\n",uniq)
```

```
    gain = entropy(examples)
```

```
    #print ("\n",gain)
```

```
    for u in uniq:
```

```
        subdata = examples[examples[attr] == u]
```

```
        #print ("\n",subdata)
```

```
        sub_e = entropy(subdata)
```

```
        gain -= (float(len(subdata)) / float(len(examples))) * sub_e
```

```
        #print ("\n",gain)
```

```
    return gain
```

Define a function named ID3 to get the decision tree for the given dataset

```
def ID3(examples, attrs):
```

```
    root = Node()
```

```
    max_gain = 0
```

```
    max_feat = ""
```

```
    for feature in attrs:
```

```

    #print ("\n",examples)
    gain = info_gain(examples, feature)
    if gain > max_gain:
        max_gain = gain
        max_feat = feature
    root.value = max_feat
    #print ("\nMax feature attr",max_feat)
    uniq = np.unique(examples[max_feat])
    #print ("\n",uniq)
    for u in uniq:
        #print ("\n",u)
        subdata = examples[examples[max_feat] == u]
        #print ("\n",subdata)
        if entropy(subdata) == 0.0:
            newNode = Node()
            newNode.isLeaf = True
            newNode.value = u
            newNode.pred = np.unique(subdata["answer"])
            root.children.append(newNode)
        else:
            dummyNode = Node()
            dummyNode.value = u
            new_attrs = attrs.copy()
            new_attrs.remove(max_feat)
            child = ID3(subdata, new_attrs)
            dummyNode.children.append(child)
            root.children.append(dummyNode)

    return root

# Define a function named printTree to draw the decision tree
def printTree(root: Node, depth=0):
    for i in range(depth):
        print("\t", end="")
    print(root.value, end="")
    if root.isLeaf:
        print(" -> ", root.pred)
    print()
    for child in root.children:
        printTree(child, depth + 1)

# Define a function named classify to classify the new example
def classify(root: Node, new):
    for child in root.children:
        if child.value == new[root.value]:
            if child.isLeaf:
                print ("Predicted Label for new example", new," is:", child.pred)
                exit
            else:
                classify (child.children[0], new)

# Finally, call the ID3, printTree and classify functions
root = ID3(data, features)
print("Decision Tree is:")
printTree(root)

```

```
print ("-----")
```

```
new = {"outlook":"sunny", "temperature":"hot", "humidity":"normal", "wind":"strong"}  
classify (root, new)
```

Output:

```
Decision Tree is:  
outlook  
  overcast -> ['yes']  
  rain  
    wind  
      strong -> ['no']  
      weak -> ['yes']  
  sunny  
    humidity  
      high -> ['no']  
      normal -> ['yes']  
-----  
Predicted Label for new example {'outlook': 'sunny', 'temperature': 'hot', 'humidity': 'normal', 'wind': 'strong'} is: ['yes']  
>
```

4. Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.

```
"""
```

```
Created on Fri Jan 27 08:29:18 2023
```

```
@author: GEOMOL GEORGE
```

```
"""
```

```
import numpy as np
```

```
X = np.array([[2, 9], [1, 5], [3, 6]], dtype=float)
```

```
y = np.array([92, 86, 89]), dtype=float)
```

```
X = X/np.amax(X,axis=0) #maximum of X array longitudinally
```

```
y = y/100
```

```
#Sigmoid Function
```

```
def sigmoid (x):
```

```
    return 1/(1 + np.exp(-x))
```

```
#Derivative of Sigmoid Function
```

```
def derivatives_sigmoid(x):
```

```
    return x * (1 - x)
```

```
#Variable initialization
```

```
epoch=5 #Setting training iterations
```

```
lr=0.1 #Setting learning rate
```

```
inputlayer_neurons = 2 #number of features in data set
```

```
hiddenlayer_neurons = 3 #number of hidden layers neurons
```

```
output_neurons = 1 #number of neurons at output layer
```

```
#weight and bias initialization
```

```
wh=np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
```

```
bh=np.random.uniform(size=(1,hiddenlayer_neurons))
```

```

wout=np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
bout=np.random.uniform(size=(1,output_neurons))

#draws a random range of numbers uniformly of dim x*y
for i in range(epoch):
    #Forward Propogation
    hinp1=np.dot(X,wh)
    hinp=hinp1 + bh
    hlayer_act = sigmoid(hinp)
    outinp1=np.dot(hlayer_act,wout)
    outinp= outinp1+bout
    output = sigmoid(outinp)

    #Backpropagation
    EO = y-output
    outgrad = derivatives_sigmoid(output)
    d_output = EO * outgrad
    EH = d_output.dot(wout.T)
    hiddengrad = derivatives_sigmoid(hlayer_act)#how much hidden layer wts contributed to error
    d_hiddenlayer = EH * hiddengrad

    wout += hlayer_act.T.dot(d_output) *lr # dotproduct of nextlayererror and currentlayerop
    wh += X.T.dot(d_hiddenlayer) *lr

    print ("-----Epoch-", i+1, "Starts-----")
    print("Input: \n" + str(X))
    print("Actual Output: \n" + str(y))
    print("Predicted Output: \n" ,output)
    print ("-----Epoch-", i+1, "Ends-----\n")

print("Input: \n" + str(X))
print("Actual Output: \n" + str(y))
print("Predicted Output: \n" ,output)

```

output:

-----Epoch- 1 Starts-----

Input:

```

[[0.66666667 1.    ]
 [0.33333333 0.55555556]
 [1.    0.66666667]]

```

Actual Output:

```

[[0.92]
 [0.86]
 [0.89]]

```

Predicted Output:

```

[[0.93250493]
 [0.92457935]
 [0.93337352]]

```

-----Epoch- 1 Ends-----

-----Epoch- 2 Starts-----

Input:

[[0.66666667 1.]
[0.33333333 0.55555556]
[1. 0.66666667]]

Actual Output:

[[0.92]
[0.86]
[0.89]]

Predicted Output:

[[0.93241372]
[0.92448608]
[0.93328291]]

-----Epoch- 2 Ends-----

-----Epoch- 3 Starts-----

Input:

[[0.66666667 1.]
[0.33333333 0.55555556]
[1. 0.66666667]]

Actual Output:

[[0.92]
[0.86]
[0.89]]

Predicted Output:

[[0.9323225]
[0.92439282]
[0.93319229]]

-----Epoch- 3 Ends-----

-----Epoch- 4 Starts-----

Input:

[[0.66666667 1.]
[0.33333333 0.55555556]
[1. 0.66666667]]

Actual Output:

[[0.92]
[0.86]
[0.89]]

Predicted Output:

[[0.93223127]
[0.92429955]
[0.93310166]]

-----Epoch- 4 Ends-----

-----Epoch- 5 Starts-----

Input:

[[0.66666667 1.]

[0.33333333 0.55555556]

[1. 0.66666667]]

Actual Output:

[[0.92]

[0.86]

[0.89]]

Predicted Output:

[[0.93214003]

[0.92420629]

[0.93301102]]

-----Epoch- 5 Ends-----

Input:

[[0.66666667 1.]

[0.33333333 0.55555556]

[1. 0.66666667]]

Actual Output:

[[0.92]

[0.86]

[0.89]]

Predicted Output:

[[0.93214003]

[0.92420629]

[0.93301102]]

5. Write a program for Implementation of K-Nearest Neighbors (K-NN) in Python

```
"""
```

Created on Fri Jan 27 09:32:02 2023

KNN

@author: GEOMOL GEORGE

```
"""
```

```
import numpy as np
```

```
import pandas as pd
```

```
dataset = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/knn1 (1).csv")
```

```
"""
```

The breast cancer dataset has the following features: Sample code number, Clump Thickness, Uniformity of Cell Size, Uniformity of Cell Shape, Marginal Adhesion, Single Epithelial Cell Size, Bare Nuclei, Bland Chromatin, Normal Nucleoli, Mitosis, Class.

```
"""
```

```
X = dataset.iloc[:, :-1].values
```

```
y = dataset.iloc[:, -1].values
```

```
dataset.shape
```

```
#splitting the dataset into the Training set and Test set
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state = 42)
```

```
#Feature Scaling
```

```
"""
```

Feature scaling is the process of converting the data into a given range.

In this case, the standard scalar technique is used.

```
"""
```

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
#Training the K-Nearest Neighbors (K-NN) Classification model on the Training set
"""
```

Once the dataset is scaled, next, the K-Nearest Neighbors (K-NN) classifier algorithm is used to create a model.

The hyperparameters such as `n_neighbors`, `metric`, and `p` are set to 5, Minkowski, and 2 respectively. The remaining hyperparameters are set to default values.

```
"""
```

```
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)
```

```
classifier.fit(X_train, y_train)
```

```
"""
```

Display the results (confusion matrix and accuracy)

Here evaluation metrics such as confusion matrix and accuracy are used to evaluate the performance of the model built using a decision tree classifier.

```
"""
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
y_pred = classifier.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```

```
=====
[[ 78   1]
 [  5  53]]
>
```

6. Write a program to implement Naïve Bayes algorithm in python and to display the results using confusion matrix and accuracy. Java/Python ML library classes can be used for this problem.

```
# -*- coding: utf-8 -*-
```

```
"""
```

Created on Fri Jan 27 11:28:03 2023

@author: GEOMOL GEORGE

Implementation of Naive Bayes in Python – Machine Learning

In this tutorial, we will understand the Implementation of Naive Bayes in Python – Machine Learning.

Importing the Necessary libraries

To begin the implementation first we will import the necessary libraries like NumPy for numerical computation and pandas for reading the dataset.

```
"""
```

```
import numpy as np
import pandas as pd
#Importing the dataset
"""
```

Next, we import or read the dataset. [Click here to download the breast cancer dataset used in this implementation.](#)

After reading the dataset, divide the dataset into concepts and targets. Store the concepts into X and

targets into y.

```
"""
```

```
dataset = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/breastcancer.csv")
```

```
X = dataset.iloc[:, :-1].values
```

```
y = dataset.iloc[:, -1].values
```

```
"""
```

Splitting the dataset into the Training set and Test set

Once the dataset is read into the memory, next, divide the dataset into two parts, training and testing using the train_test_split function from sklearn.

The test_size and random_state attributes are set to 0.25 and 0 respectively.

You can change these attributes as per your requirements.

```
"""
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)
```

#Feature Scaling

```
"""
```

Feature scaling is the process of converting the data into a min-max range. In this case, the standard scalar method is used.

```
"""
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
X_train = sc.fit_transform(X_train)
```

```
X_test = sc.transform(X_test)
```

```
"""
```

Training the Naive Bayes Classification model on the Training set

Once the dataset is scaled, next, the Naive Bayes classifier algorithm is used to create a model.

The GaussianNB function is imported from sklearn.naive_bayes library. The hyperparameters such as kernel, and random_state to linear, and 0 respectively. The remaining hyperparameters of the support vector machine algorithm are set to default values.

```
"""
```

```
from sklearn.naive_bayes import GaussianNB
```

```
classifier = GaussianNB()
```

```
classifier.fit(X_train, y_train)
```

#Naive Bayes classifier model

```
GaussianNB(priors=None, var_smoothing=1e-09)
```

#Display the results (confusion matrix and accuracy)

```
"""
```

Here evaluation metrics such as confusion matrix and accuracy are used to evaluate the performance of the model built using a decision tree classifier.

```
"""
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
```

```
y_pred = classifier.predict(X_test)
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
print(cm)
```

```
accuracy_score(y_test, y_pred)
```

output:

```
===== RESTART
```

```
[[ 99  8]
 [ 2 62]]
```

```
>>>
```

7. Write a program to implement Linear Regression (LR) algorithm in python

```
"""
```

Created on Fri Jan 27 12:10:19 2023

@author: GEOMOL GEORGE

''''''

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
dataset = pd.read_csv('D:/GEO/BE COURSES/LAB/DATASET/Salary_Data.csv')
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
```

```
dataset.head()
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 1/3, random_state = 0)
```

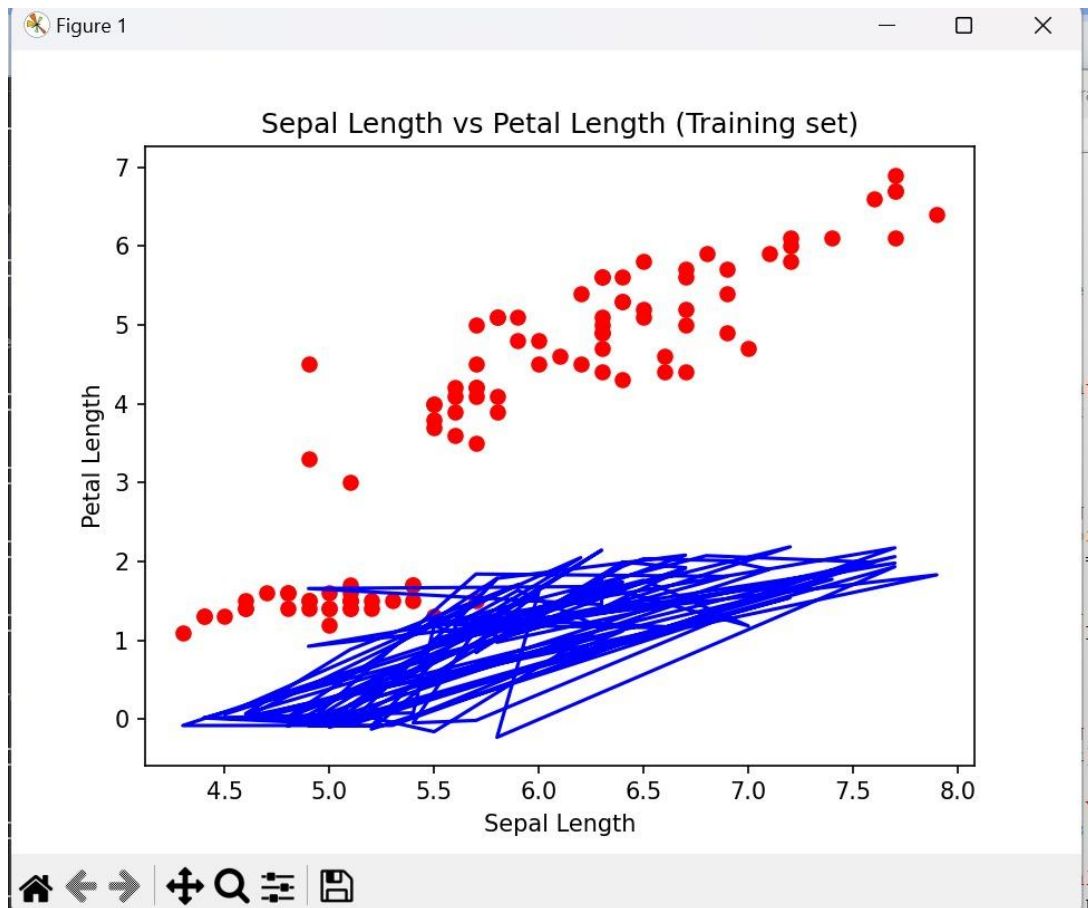
```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
y_pred = regressor.predict(X_test)
pd.DataFrame(data={'Actuals': y_test, 'Predictions': y_pred})
```

#Visualising the Training set results Here scatter plot is used to visualize the results.

```
plt.scatter(X_train, y_train, color = 'red')
plt.plot(X_train, regressor.predict(X_train), color = 'blue')
plt.title('Salary vs Experience (Training set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.show()
```



8. Implementation Of Linear And Polynomial Regression In Python

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
dataset = pd.read_csv('D:/GEO/BE COURSES/LAB/DATASET/Position_Salaries.csv')
X = dataset.iloc[:, 1:-1].values
y = dataset.iloc[:, -1].values
"""
```

Training the Linear Regression model on the Whole dataset

A Linear regression algorithm is used to create a model.

A LinearRegression function is imported from sklearn.linear_model library.

"""

```
from sklearn.linear_model import LinearRegression
```

```
lin_reg = LinearRegression()
```

```
lin_reg.fit(X, y)
```

#Linear Regression classifier model

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

"""

Training the Polynomial Regression model on the Whole dataset

A polynomial regression algorithm is used to create a model.

"""

```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly_reg = PolynomialFeatures(degree = 4)
```

```
X_poly = poly_reg.fit_transform(X)
```

```
lin_reg_2 = LinearRegression()
```

```
lin_reg_2.fit(X_poly, y)
```

```
#Polynomial Regression classifier model
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
"""
```

```
Visualising the Linear Regression results
```

Here scatter plot is used to visualize the results. The title of the plot is set to Truth or Bluff (Linear Regression), xlabel is set to Position Level , and ylabel is set to Salary.

```
"""
```

```
plt.scatter(X, y, color = 'red')
```

```
plt.plot(X, lin_reg.predict(X), color = 'blue')
```

```
plt.title('Truth or Bluff (Linear Regression)')
```

```
plt.xlabel('Position Level')
```

```
plt.ylabel('Salary')
```

```
plt.show()
```

```
#Visualising the Polynomial Regression results
```

```
"""
```

The title of the plot is set to Truth or Bluff (Polynomial Regression), xlabel is set to Position level, and ylabel is set to Salary.

```
"""
```

```
plt.scatter(X, y, color = 'red')
```

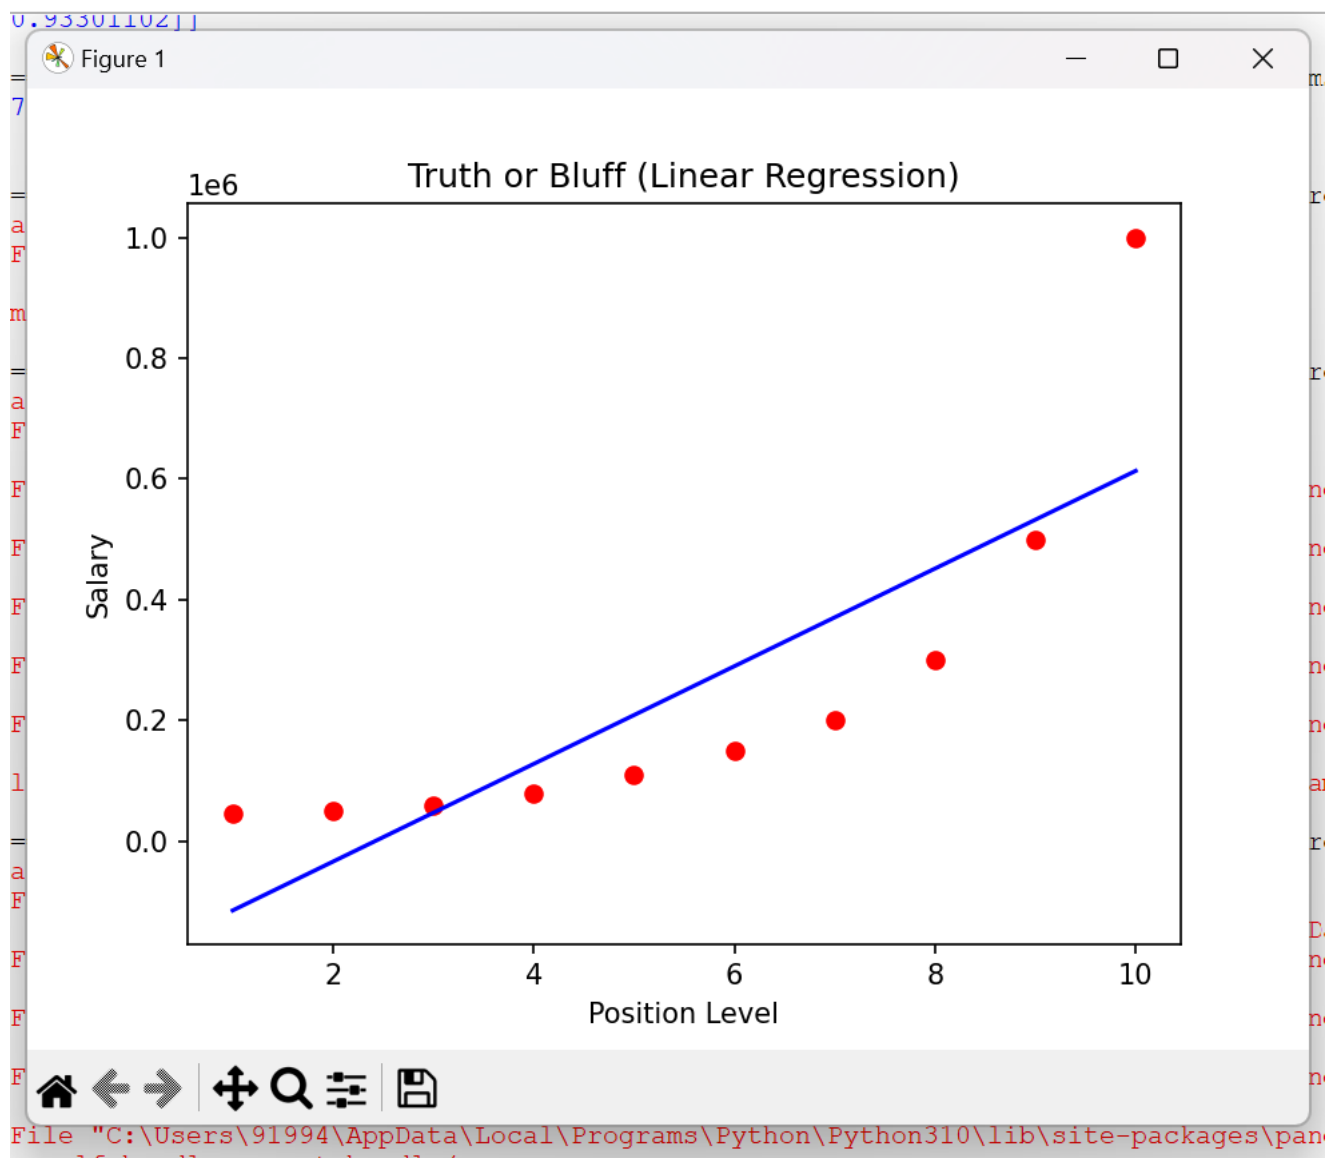
```
plt.plot(X, lin_reg_2.predict(poly_reg.fit_transform(X)), color = 'blue')
```

```
plt.title('Truth or Bluff (Polynomial Regression)')
```

```
plt.xlabel('Position level')
```

```
plt.ylabel('Salary')
```

```
plt.show()
```



9. Write a program to implement Logistic Regression (LR) algorithm in python

.....

Created on Fri Jan 27 11:53:58 2023

@author: GEOMOL GEORGE

.....

.....

Implementation of Logistic Regression (LR) in Python – Machine Learning

In this tutorial, we will understand the Implementation of Logistic Regression (LR) in Python – Machine Learning.

Importing the libraries

.....

import numpy as np

import pandas as pd

#"Importing the dataset

.....

After importing the necessary libraries, next, we import or read the dataset.

[Click here to download the breast cancer dataset used in this implementation.](#)

The breast cancer dataset has the following features:

Sample code number, Clump Thickness, Uniformity of Cell Size, Uniformity of Cell Shape, Marginal Adhesion,

Single Epithelial Cell Size, Bare Nuclei, Bland Chromatin, Normal Nucleoli, Mitosis, Class.

"""

divide the dataset into concepts and targets. Store the concepts into X and targets into y.

```
dataset = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/breastcancer.csv")
```

```
X = dataset.iloc[:, :-1].values
```

```
y = dataset.iloc[:, -1].values
```

#Splitting the dataset into the Training set and Test

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_state = 2)
```

#Feature Scaling

"""

Feature scaling is the process of converting the data into a given range. In this case, the standard scalar technique is used.

```
from sklearn.preprocessing import StandardScaler
```

"""

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
X_train = sc.fit_transform(X_train)
```

```
X_test = sc.transform(X_test)
```

"""

Training the Logistic Regression (LR) Classification model on the Training set

Once the dataset is scaled, next, the Logistic Regression (LR) classifier algorithm is used to create a model.

The hyperparameters such as random_state to 0 respectively.

The remaining hyperparameters Logistic Regression (LR) are set to default values.

"""

```
from sklearn.linear_model import LogisticRegression
```

```
classifier = LogisticRegression(random_state = 0)
```

```
classifier.fit(X_train, y_train)
```

#Logistic Regression (LR) classifier model

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
```

```
    intercept_scaling=1, l1_ratio=None, max_iter=100,
```

```
    multi_class='warn', n_jobs=None, penalty='l2',
```

```
    random_state=0, solver='warn', tol=0.0001, verbose=0,
```

```
    warm_start=False)
```

#Display the results (confusion matrix and accuracy)

"""

Here evaluation metrics such as confusion matrix and accuracy are used to evaluate the performance of the model

built using a decision tree classifier.

"""

```
from sklearn.metrics import confusion_matrix, accuracy_score
```

```
y_pred = classifier.predict(X_test)
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
print(cm)
```

```
accuracy_score(y_test, y_pred)
```



```
=====
[[117  8]
 [ 6 74]]
Accuracy: 0.9317073170731708
>>
```

10. Python Program to Implement Estimation & MAXimization Algorithm Program:

"""

Created on Fri Jan 27 13:13:36 2023

@author: GEOMOL GEORGE

Python Program to Implement Estimation & MAXimization Algorithm

"""

```
#from sklearn.cluster import KMeans
```

```
from sklearn.mixture import GaussianMixture
```

```
import sklearn.metrics as metrics
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
names = ['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width', 'Class']
```

```
dataset = pd.read_csv('D:/GEO/BE COURSES/LAB/DATASET/8-dataset.csv', names=names)
```

```
X = dataset.iloc[:, :-1]
```

```
label = {'Iris-setosa': 0, 'Iris-versicolor': 1, 'Iris-virginica': 2}
```

```
y = [label[c] for c in dataset.iloc[:, -1]]
```

```

plt.figure(figsize=(14,7))

colormap=np.array(['red','lime','black'])

# REAL PLOT

plt.subplot(1,3,1)

plt.title('Real')

plt.scatter(X.Petal_Length,X.Petal_Width,c=colormap[y])


# GMM PLOT

gmm=GaussianMixture(n_components=3, random_state=0).fit(X)

y_cluster_gmm=gmm.predict(X)

plt.subplot(1,3,3)

plt.title('GMM Classification')

plt.scatter(X.Petal_Length,X.Petal_Width,c=colormap[y_cluster_gmm])


print('The accuracy score of EM: ',metrics.accuracy_score(y, y_cluster_gmm))

print('The Confusion matrix of EM:\n ',metrics.confusion_matrix(y, y_cluster_gmm))

```

OUTPUT:

```

> The accuracy score of EM:  0.36666666666666664

```

- 11. There are three credit scores that banks and credit card companies use to label their customers: Good, Standard, Poor A person with a good credit score will get loans from any bank and financial institution. Write a program for the task of Credit Score Classification, we need a labelled dataset with credit scores.**

"""

Credit Score Classification

There are three credit scores that banks and credit card companies use to label their customers:

Good

Standard

Poor

A person with a good credit score will get loans from any bank and financial institution. For the task of Credit Score Classification, we need a labelled dataset with credit scores.

```
"""
```

```
import pandas as pd
```

```
import numpy as np
```

```
import plotly.express as px
```

```
import plotly.graph_objects as go
```

```
import plotly.io as pio
```

```
pio.templates.default = "plotly_white"
```

```
import plotly.io as io
```

```
io.renderers.default='browser'
```

```
data = pd.read_csv('D:/GEO/BE COURSES/LAB/DATASET/Credit-Score-Data/Credit Score  
Data/train.csv')
```

```
print(data.head())
```

```
print(data.info())
```

```
#the dataset has any null values or not:
```

```
print(data.isnull().sum())
```

```
#The dataset doesn't have any null values. As this dataset is labelled, let's have a look at the  
Credit_Score column values:
```

```
data["Credit_Score"].value_counts()
```

```
data.shape
```

```
#Data Exploration
```

```
"""
```

The dataset has many features that can train a Machine Learning model for credit score classification.

Let's explore all the features one by one.

I will start by exploring the occupation feature to know if the occupation of the person affects credit scores:

```
"""  
  
fig = px.box(data,  
  
              x="Occupation",  
  
              color="Credit_Score",  
  
              title="Credit Scores Based on Occupation",  
  
              color_discrete_map={'Poor':'red',  
                                  'Standard':'yellow',  
                                  'Good':'green'})  
  
fig.show()  
"""
```

There's not much difference in the credit scores of all occupations mentioned in the data. Now let's explore

whether the Annual Income of the person impacts your credit scores or not:

```
"""  
  
fig = px.box(data,  
  
              x="Credit_Score",  
  
              y="Annual_Income",  
  
              color="Credit_Score",  
  
              title="Credit Scores Based on Annual Income",  
  
              color_discrete_map={'Poor':'red',  
                                  'Standard':'yellow',  
                                  'Good':'green'})  
  
fig.update_traces(quartilemethod="exclusive")  
  
fig.show()
```

```
"""
```

let's explore whether the monthly in-hand salary impacts credit scores or not:

```
"""
```

```
fig = px.box(data,  
             x="Credit_Score",  
             y="Monthly_Inhand_Salary",  
             color="Credit_Score",  
             title="Credit Scores Based on Monthly Inhand Salary",  
             color_discrete_map={'Poor':'red',  
                                 'Standard':'yellow',  
                                 'Good':'green'})  
  
fig.update_traces(quartilemethod="exclusive")  
  
fig.show()
```

```
fig = px.box(data,  
             x="Credit_Score",  
             y="Num_Bank_Accounts",  
             color="Credit_Score",  
             title="Credit Scores Based on Number of Bank Accounts",  
             color_discrete_map={'Poor':'red',  
                                 'Standard':'yellow',  
                                 'Good':'green'})  
  
fig.update_traces(quartilemethod="exclusive")  
  
fig.show()
```

impact on credit scores based on the number of credit cards you have:

```
fig = px.box(data,  
             x="Credit_Score",
```

```

y="Num_Credit_Card",

color="Credit_Score",

title="Credit Scores Based on Number of Credit cards",

color_discrete_map={'Poor':'red',

                    'Standard':'yellow',

                    'Good':'green'})

fig.update_traces(quartilemethod="exclusive")

fig.show()

```

```

fig = px.box(data,

             x="Credit_Score",

             y="Interest_Rate",

             color="Credit_Score",

             title="Credit Scores Based on the Average Interest rates",

             color_discrete_map={'Poor':'red',

                                 'Standard':'yellow',

                                 'Good':'green'})

fig.update_traces(quartilemethod="exclusive")

fig.show()

```

```

data["Credit_Mix"] = data["Credit_Mix"].map({"Standard": 1,

                                             "Good": 2,

                                             "Bad": 0})

from sklearn.model_selection import train_test_split

x = np.array(data[["Annual_Income", "Monthly_Inhand_Salary",

                  "Num_Bank_Accounts", "Num_Credit_Card",

                  "Interest_Rate", "Num_of_Loan",

                  "Delay_from_due_date", "Num_of_Delayed_Payment",

```

```

        "Credit_Mix", "Outstanding_Debt",

        "Credit_History_Age", "Monthly_Balance"]])

y = np.array(data[["Credit_Score"]])

xtrain, xtest, ytrain, ytest = train_test_split(x, y,

                                                test_size=0.33,

                                                random_state=42)

from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier()

model.fit(xtrain, ytrain)

print("Credit Score Prediction : ")

a = float(input("Annual Income: "))

b = float(input("Monthly Inhand Salary: "))

c = float(input("Number of Bank Accounts: "))

d = float(input("Number of Credit cards: "))

e = float(input("Interest rate: "))

f = float(input("Number of Loans: "))

g = float(input("Average number of days delayed by the person: "))

h = float(input("Number of delayed payments: "))

i = input("Credit Mix (Bad: 0, Standard: 1, Good: 3) : ")

j = float(input("Outstanding Debt: "))

k = float(input("Credit History Age: "))

l = float(input("Monthly Balance: "))

features = np.array([[a, b, c, d, e, f, g, h, i, j, k, l]])

print("Predicted Credit Score = ", model.predict(features))

```

Output:

```
ID Customer_ID ... Monthly_Balance Credit_Score
```

| | | | | |
|---|------|----------|------------|------|
| 0 | 5634 | 3392 ... | 312.494089 | Good |
| 1 | 5635 | 3392 ... | 284.629162 | Good |
| 2 | 5636 | 3392 ... | 331.209863 | Good |
| 3 | 5637 | 3392 ... | 223.451310 | Good |
| 4 | 5638 | 3392 ... | 341.489231 | Good |

[5 rows x 28 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 100000 entries, 0 to 99999

Data columns (total 28 columns):

| # | Column | Non-Null Count | Dtype |
|----|------------------------|-----------------|---------|
| 0 | ID | 100000 non-null | int64 |
| 1 | Customer_ID | 100000 non-null | int64 |
| 2 | Month | 100000 non-null | int64 |
| 3 | Name | 100000 non-null | object |
| 4 | Age | 100000 non-null | float64 |
| 5 | SSN | 100000 non-null | float64 |
| 6 | Occupation | 100000 non-null | object |
| 7 | Annual_Income | 100000 non-null | float64 |
| 8 | Monthly_Inhand_Salary | 100000 non-null | float64 |
| 9 | Num_Bank_Accounts | 100000 non-null | float64 |
| 10 | Num_Credit_Card | 100000 non-null | float64 |
| 11 | Interest_Rate | 100000 non-null | float64 |
| 12 | Num_of_Loan | 100000 non-null | float64 |
| 13 | Type_of_Loan | 100000 non-null | object |
| 14 | Delay_from_due_date | 100000 non-null | float64 |
| 15 | Num_of_Delayed_Payment | 100000 non-null | float64 |

16 Changed_Credit_Limit 100000 non-null float64
17 Num_Credit_Inquiries 100000 non-null float64
18 Credit_Mix 100000 non-null object
19 Outstanding_Debt 100000 non-null float64
20 Credit_Utilization_Ratio 100000 non-null float64
21 Credit_History_Age 100000 non-null float64
22 Payment_of_Min_Amount 100000 non-null object
23 Total_EMI_per_month 100000 non-null float64
24 Amount_invested_monthly 100000 non-null float64
25 Payment_Behaviour 100000 non-null object
26 Monthly_Balance 100000 non-null float64
27 Credit_Score 100000 non-null object

dtypes: float64(18), int64(3), object(7)

memory usage: 21.4+ MB

None

| | |
|-----------------------|---|
| ID | 0 |
| Customer_ID | 0 |
| Month | 0 |
| Name | 0 |
| Age | 0 |
| SSN | 0 |
| Occupation | 0 |
| Annual_Income | 0 |
| Monthly_Inhand_Salary | 0 |
| Num_Bank_Accounts | 0 |
| Num_Credit_Card | 0 |
| Interest_Rate | 0 |
| Num_of_Loan | 0 |

```
Type_of_Loan          0
Delay_from_due_date   0
Num_of_Delayed_Payment 0
Changed_Credit_Limit  0
Num_Credit_Inquiries  0
Credit_Mix           0
Outstanding_Debt      0
Credit_Utilization_Ratio 0
Credit_History_Age    0
Payment_of_Min_Amount 0
Total_EMI_per_month   0
Amount_invested_monthly 0
Payment_Behaviour     0
Monthly_Balance       0
Credit_Score          0
dtype: int64
```

Warning (from warnings module):

File "C:\machine learning programs\lab\11th.py", line 111

```
model.fit(xtrain, ytrain)
```

DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

Credit Score Prediction :

Annual Income: 50000

Monthly Inhand Salary: 2000

Number of Bank Accounts: 3

Number of Credit cards: 2

Interest rate: 3.5

Number of Loans: 0

Average number of days delayed by the person: 0

Number of delayed payments: 0

Credit Mix (Bad: 0, Standard: 1, Good: 3) : 23

Outstanding Debt: 2

Credit History Age: 4

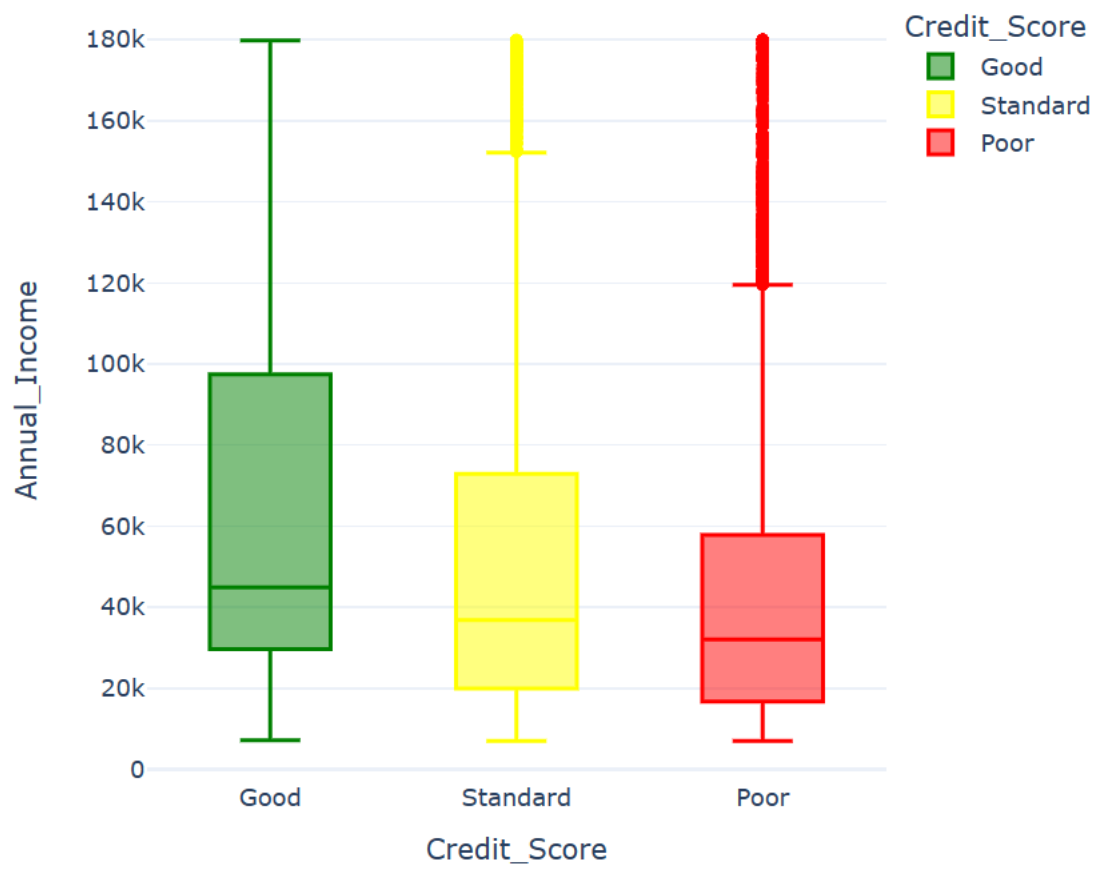
Monthly Balance: 1600

Predicted Credit Score = ['Good']


Graphs:



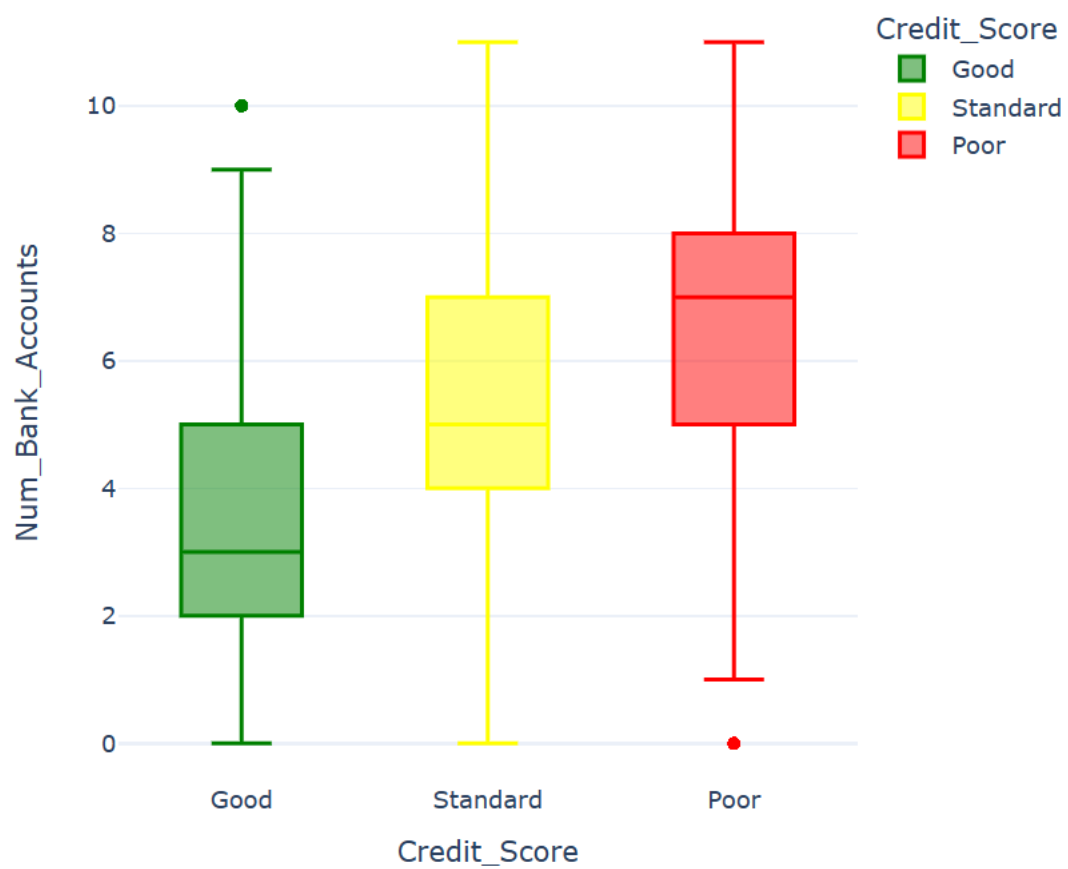
Credit Scores Based on Annual Income




Credit Scores Based on Monthly Inhand Salary




Credit Scores Based on Number of Bank Accounts



Credit Scores Based on Number of Credit cards





12. Iris Flower Classification using KNN

"""

Created on Tue Jan 24 09:37:46 2023

@author: GEOMOL GEORGE

Iris Flower Classification using Python

"""

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
iris = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/IRIS.csv")
```

```
#first five rows of this dataset:
```



```

print(iris.head())

print(iris.describe())

#The target labels of this dataset are present in the species column, let's have a quick look at the target labels:

print("Target Labels", iris["species"].unique())

#plot the data using a scatter plot which will plot the iris species according to the sepal length and sepal width:

import plotly.io as io

io.renderers.default='browser'

import plotly.express as px

fig = px.scatter(iris, x="sepal_width", y="sepal_length", color="species")

fig.show()

#Iris Classification Model

x = iris.drop("species", axis=1)

y = iris["species"]

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)

from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=1)

knn.fit(x_train, y_train)

x_new = np.array([[6, 2.9, 1, 0.2]])


prediction = knn.predict(x_new) print

("Prediction: {}".format(prediction))

```

Output:

| | sepal_length | sepal_width | petal_length | petal_width | species |
|---------------|--|-------------|--------------|-------------|-------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |
| | sepal_length | sepal_width | petal_length | petal_width | |
| count | 150.000000 | 150.000000 | 150.000000 | 150.000000 | |
| mean | 5.843333 | 3.054000 | 3.758667 | 1.198667 | |
| std | 0.828066 | 0.433594 | 1.764420 | 0.763161 | |
| min | 4.300000 | 2.000000 | 1.000000 | 0.100000 | |
| 25% | 5.100000 | 2.800000 | 1.600000 | 0.300000 | |
| 50% | 5.800000 | 3.000000 | 4.350000 | 1.300000 | |
| 75% | 6.400000 | 3.300000 | 5.100000 | 1.800000 | |
| max | 7.900000 | 4.400000 | 6.900000 | 2.500000 | |
| Target Labels | ['Iris-setosa' 'Iris-versicolor' 'Iris-virginica'] | | | | |



13. Car Price Prediction Model using Python

"""

Created on Tue Jan 24 10:02:09 2023

@author: GEOMOL GEORGE

car price prediction

"""

import numpy as np

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor

data = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/CarPrice.csv")
data.head()
data.shape
data.isnull().sum()

#So this dataset doesn't have any null values, now let's look at some of the other important insights to get
#an idea of what kind of data we're dealing with:
data.info()
data.describe()
data.CarName.unique()

sns.set_style("whitegrid")
plt.figure(figsize=(15, 10))
sns.distplot(data.price)
plt.show()

#Now let's have a look at the correlation among all the features of this dataset:

print(data.corr())

plt.figure(figsize=(20, 15))
correlations = data.corr()
sns.heatmap(correlations, cmap="coolwarm", annot=True)
plt.show()

#Training a Car Price Prediction Model
predict = "price"
data = data[["symboling", "wheelbase", "carlength",
```

```
"carwidth", "carheight", "curbweight",  
"enginesize", "boreratio", "stroke",  
"compressionratio", "horsepower", "peakrpm",  
"citympg", "highwaympg", "price"]]
```

```
x = np.array(data.drop([predict], 1))
```

```
y = np.array(data[predict])
```

```
from sklearn.model_selection import train_test_split
```

```
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.2)
```

```
from sklearn.tree import DecisionTreeRegressor
```

```
model = DecisionTreeRegressor()
```

```
model.fit(xtrain, ytrain)
```

```
predictions = model.predict(xtest)
```

```
from sklearn.metrics import mean_absolute_error
```


```
model.score(xtest, predictions)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                 205 non-null    int64
1   symboling              205 non-null    int64
2   CarName                205 non-null    object
3   fueltype               205 non-null    object
4   aspiration              205 non-null    object
5   doornumber             205 non-null    object
6   carbody                205 non-null    object
7   drivewheel             205 non-null    object
8   enginelocation         205 non-null    object
9   wheelbase              205 non-null    float64
10  carlength              205 non-null    float64
11  carwidth               205 non-null    float64
12  carheight              205 non-null    float64
13  curbweight             205 non-null    int64
14  enginetype             205 non-null    object
15  cylindernumber         205 non-null    object
16  enginesize             205 non-null    int64
17  fuelsystem             205 non-null    object
18  boreratio              205 non-null    float64
19  stroke                 205 non-null    float64
20  compressionratio       205 non-null    float64
21  horsepower             205 non-null    int64
22  peakrpm                205 non-null    int64
23  citympg                205 non-null    int64
24  highwaympg             205 non-null    int64
25  price                  205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

Figure 1



x=2.971e+04 y=8.7e-06

14. House Price Prediction using Machine Learning in Python

```
# -*- coding: utf-8 -*-
```

```
"""
```

Created on Mon Jan 23 11:39:55 2023

```
@author: GEOMOL GEORGE
```

```
"""
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
dataset = pd.read_excel("HousePricePrediction.xlsx")
```

```
# Printing first 5 records of the dataset
```

```
print(dataset.head(5))
```

```
dataset.shape
```

```
obj = (dataset.dtypes == 'object')
```

```
object_cols = list(obj[obj].index)
```

```
print("Categorical variables:",len(object_cols))
```

```
int_ = (dataset.dtypes == 'int')
```

```
num_cols = list(int_[int_].index)
```

```
print("Integer variables:",len(num_cols))
```

```
fl = (dataset.dtypes == 'float')
```

```
fl_cols = list(fl[fl].index)
```

```
print("Float variables:",len(fl_cols))
```

```
plt.figure(figsize=(12, 6))
```

```
sns.heatmap(dataset.corr(),
```

```
                cmap = 'BrBG',
```

```
                fmt = '.2f',
```

```
                linewidths = 2,
```

```
                annot = True)
```

```
unique_values = []
```

```

for col in object_cols:
    unique_values.append(dataset[col].unique().size)

plt.figure(figsize=(10,6))
plt.title('No. Unique values of Categorical Features')
plt.xticks(rotation=90)
sns.barplot(x=object_cols,y=unique_values)


plt.figure(figsize=(18, 36))
plt.title('Categorical Features: Distribution')
plt.xticks(rotation=90)
index = 1


for col in object_cols:
    y = dataset[col].value_counts()
    plt.subplot(11, 4, index)
    plt.xticks(rotation=90)
    sns.barplot(x=list(y.index), y=y)
    index += 1


dataset.drop(['Id'],
              axis=1,
              inplace=True)


dataset['SalePrice'] = dataset['SalePrice'].fillna(dataset['SalePrice'].mean())


new_dataset = dataset.dropna()


new_dataset.isnull().sum()


from sklearn.preprocessing import OneHotEncoder


s = (new_dataset.dtypes == 'object')
object_cols = list(s[s].index)
print("Categorical variables:")
print(object_cols)
print('No. of. categorical features: ',len(object_cols))

```

```

OH_encoder = OneHotEncoder(sparse=False)
OH_cols = pd.DataFrame(OH_encoder.fit_transform(new_dataset[object_cols]))
OH_cols.index = new_dataset.index
OH_cols.columns = OH_encoder.get_feature_names()
df_final = new_dataset.drop(object_cols, axis=1)
df_final = pd.concat([df_final, OH_cols], axis=1)


from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import train_test_split


X = df_final.drop(['SalePrice'], axis=1)
Y = df_final['SalePrice']


# Split the training set into
# training and validation set
X_train, X_valid, Y_train, Y_valid = train_test_split(X, Y, train_size=0.8, test_size=0.2, random_state=0)
#Model and Accuracy
#svm

from sklearn import svm
from sklearn.svm import SVC
from sklearn.metrics import mean_absolute_percentage_error


model_SVR = svm.SVR()
model_SVR.fit(X_train,Y_train)
Y_pred = model_SVR.predict(X_valid)


print(mean_absolute_percentage_error(Y_valid, Y_pred))
#Random forest

from sklearn.ensemble import RandomForestRegressor


model_RFR = RandomForestRegressor(n_estimators=10)
model_RFR.fit(X_train, Y_train)
Y_pred = model_RFR.predict(X_valid)


mean_absolute_percentage_error(Y_valid, Y_pred)

```



```
#LinearRegression

from sklearn.linear_model import LinearRegression

model_LR = LinearRegression()
model_LR.fit(X_train, Y_train)
Y_pred = model_LR.predict(X_valid)

print(mean_absolute_percentage_error(Y_valid, Y_pred))
```

```
===== RESTART: C:\machin
   Id  MSSubClass  MSZoning  ...  BsmtFinSF2  TotalBsmtSF  SalePrice
0    0           60       RL  ...           0.0           856.0    208500.0
1    1           20       RL  ...           0.0          1262.0    181500.0
2    2           60       RL  ...           0.0           920.0    223500.0
3    3           70       RL  ...           0.0           756.0    140000.0
4    4           60       RL  ...           0.0          1145.0    250000.0

[5 rows x 13 columns]
Categorical variables: 4
Integer variables: 0
Float variables: 3
```

15. NAÏVE IRIS CLASSIFICATION

-*- coding: utf-8 -*-

"""

Created on Tue Jan 24 11:44:42 2023

@author: GEOMOL GEORGE

"""

```
from sklearn.naive_bayes import GaussianNB
from sklearn.naive_bayes import MultinomialNB
from sklearn import datasets
from sklearn.metrics import confusion_matrix

iris = datasets.load_iris()
gnb = GaussianNB()
mnb = MultinomialNB()

y_pred_gnb = gnb.fit(iris.data, iris.target).predict(iris.data)

cnf_matrix_gnb = confusion_matrix(iris.target, y_pred_gnb)

print(cnf_matrix_gnb)
y_pred_mnb = mnb.fit(iris.data, iris.target).predict(iris.data)
cnf_matrix_mnb = confusion_matrix(iris.target, y_pred_mnb)
print(cnf_matrix_mnb)
```

```

=====
[[50  0  0]
 [ 0 47  3]
 [ 0  3 47]]
[[50  0  0]
 [ 0 46  4]
 [ 0  3 47]]
>

```

16. Comparison of Classification Algorithms

"""

Created on Wed Jan 25 13:10:08 2023

@author: GEOMOL GEORGE

Comparison of Classification Algorithms

"""

```

import numpy
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import BernoulliNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import PassiveAggressiveClassifier
from sklearn.metrics import classification_report

iris= pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/IRIS.csv")
print(iris.head())

x = iris.drop("species", axis=1)
y = iris["species"]
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.10, random_state=42)

#x = np.array(data[["Age", "EstimatedSalary"]])
#y = np.array(data[["Purchased"]])

#xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.10, random_state=42)
decisiontree = DecisionTreeClassifier()
logisticregression = LogisticRegression()
knearestclassifier = KNeighborsClassifier()
#svm_classifier = SVC()
bernoulli_naiveBayes = BernoulliNB()
passiveAggressive = PassiveAggressiveClassifier()

knearestclassifier.fit(x_train, y_train)
decisiontree.fit(x_train, y_train)
logisticregression.fit(x_train, y_train)
passiveAggressive.fit(x_train, y_train)

data1 = {"Classification Algorithms": ["KNN Classifier", "Decision Tree Classifier",
                                     "Logistic Regression", "Passive Aggressive Classifier"],
        "Score": [knearestclassifier.score(x,y), decisiontree.score(x, y),
                  logisticregression.score(x, y), passiveAggressive.score(x,y)] }
score = pd.DataFrame(data1)

```

score

```
>> ===== RESTART: C:\machine le
      sepal_length  sepal_width  petal_length  petal_width  species
0           5.1         3.5         1.4         0.2  Iris-setosa
1           4.9         3.0         1.4         0.2  Iris-setosa
2           4.7         3.2         1.3         0.2  Iris-setosa
3           4.6         3.1         1.5         0.2  Iris-setosa
4           5.0         3.6         1.4         0.2  Iris-setosa
Warning (from warnings module):
```

17.Mobile Price Classification using Python

"""

Created on Wed Jan 25 13:32:56 2023

@author: GEOMOL GEORGE

Mobile Price Classification using Python

"""

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score

sns.set()

import plotly.io as io

io.renderers.default='browser'

data = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/mobile_prices.csv")

print(data.head())

plt.figure(figsize=(12, 10))

sns.heatmap(data.corr(), annot=True, cmap="coolwarm", linecolor='white', linewidths=1)

#data preparation

x = data.iloc[:, :-1].values

y = data.iloc[:, -1].values

x = StandardScaler().fit_transform(x)

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=0)
```

```
# Logistic Regression algorithm provided by Scikit-learn:
```

```
from sklearn.linear_model import LogisticRegression
```

```
lreg = LogisticRegression()
```

```
lreg.fit(x_train, y_train)
```

```
y_pred = lreg.predict(x_test)
```

```
#accuracy of the model:
```

```
accuracy = accuracy_score(y_test, y_pred) * 100
```

```
print("Accuracy of the Logistic Regression Model: ",accuracy)
```

```
#predictions made by the model:
```

```
print(y_pred)
```

```
#Let's have a look at the number of mobile phones classified for each price range:
```

```
(unique, counts) = np.unique(y_pred, return_counts=True)
```

```
price_range = np.asarray((unique, counts)).T
```

```
print(price_range)
```

```
>>> ===== RESTART: C:\machine learning programs\lab\17th.py =====
  battery_power  blue  clock_speed  ...  touch_screen  wifi  price_range
0             842     0           2.2  ...           0       1             1
1            1021     1           0.5  ...           1       0             2
2             563     1           0.5  ...           1       0             2
3             615     1           2.5  ...           0       0             2
4            1821     1           1.2  ...           1       0             1

[[5 rows x 21 columns]
Accuracy of the Logistic Regression Model: 95.5
[[3 0 2 2 3 0 0 3 3 1 1 3 0 2 3 0 3 2 2 1 0 0 3 1 2 2 3 1 3 1 1 0 2 0 2 3 0
 0 3 3 3 1 3 3 1 3 0 1 3 1 1 3 0 3 0 2 2 2 0 3 3 1 3 2 1 2 3 2 2 2 3 2 1 0
 1 3 2 2 1 2 3 3 3 0 0 0 2 1 2 3 1 2 2 1 0 3 3 3 0 3 1 1 3 1 3 2 2 3 2 3 3
 0 0 1 3 3 0 0 1 0 0 3 2 2 1 2 1 1 0 2 1 3 3 3 3 3 2 0 1 1 2 1 3 0 3 0 0
 2 0 1 1 1 1 3 0 0 3 1 3 2 1 3 1 2 3 3 2 1 0 3 1 2 3 3 0 2 3 1 2 1 0 1 2
 2 2 0 3 3 1 1 0 2 3 0 1 2 2 0 3 3 3 1 2 3 3 3 0 0 2 3 3 0 0 1 3 2 3 3 3
 0 0 2 3 3 1 0 2 0 0 0 3 2 1 2 2 1 1 0 2 3 3 0 0 1 3 3 1 3 0 3 1 1 0 2 3 3
 2 0 0 1 2 3 2 2 3 2 1 0 3 3 2 1 3 2 2 2 1 0 2 2 1 0 0 2 2 2 3 0 1 3 0 2 2
 3 0 2 0 1 1 3 0 0 2 3 1 2 0 2 0 3 0 3 3 2 3 1 2 2 1 1 1 0 1 0 3 1 0 3 0 0
 1 3 0 3 1 1 0 1 3 0 2 1 1 2 1 1 0 2 0 0 3 1 2 3 2 2 0 3 2 2 1 3 2 3 3 0
 2 0 3 0 1 1 2 3 1 3 1 2 0 1 2 3 0 0 1 3 0 3 0 2 2 1 1 0 2 0]]
[[ 0 95]
 [ 1 90]
 [ 2 97]
 [ 3 118]]
```

18. Perceptron IRIS classification

```
"""
```

Created on Tue Jan 24 10:51:21 2023

Perceptron IRIS classification

@author: GEOMOL GEORGE

''''

```
from sklearn import datasets
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Perceptron
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

iris = datasets.load_iris()
X = iris.data[:, [2, 3]]
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=1, stratify=y)

sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)

ppn = Perceptron(eta0=0.1, random_state=1)
ppn.fit(X_train_std, y_train)

y_pred = ppn.predict(X_test_std)

print('Accuracy: %.3f' % accuracy_score(y_test, y_pred))
print('Accuracy: %.3f' % ppn.score(X_test_std, y_test))
```

```
... | =====
    | Accuracy: 0.978
    | Accuracy: 0.978
>> |
```

19.Implementation of Naive Bayes in Python – Machine Learning

''''

Created on Fri Jan 27 11:28:03 2023

@author: GEOMOL GEORGE

Implementation of Naive Bayes in Python – Machine Learning

In this tutorial, we will understand the Implementation of Naive Bayes in Python – Machine Learning.

Importing the Necessary libraries

```
"""
```

```
import numpy as np
```

```
import pandas as pd
```

```
#Importing the dataset
```

```
"""
```

Next, we import or read the dataset. [Click here to download the breast cancer dataset used in this implementation](#)

```
"""
```

```
dataset = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/breastcancer.csv")
```

```
X = dataset.iloc[:, :-1].values
```

```
y = dataset.iloc[:, -1].values
```

```
"""
```

Splitting the dataset into the Training set and Test set

Once the dataset is read into the memory, next, divide the dataset into two parts, training and testing using the `train_test_split` function from `sklearn`.

The `test_size` and `random_state` attributes are set to 0.25 and 0 respectively.

You can change these attributes as per your requirements.

```
"""
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)
```

```
#Feature Scaling
```

```
"""
```

Feature scaling is the process of converting the data into a min-max range. In this case, the standard scalar method is used.

```
"""
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
X_train = sc.fit_transform(X_train)
```

```
X_test = sc.transform(X_test)
```

```
"""
```

Training the Naive Bayes Classification model on the Training set

The `GaussianNB` function is imported from `sklearn.naive_bayes` library. The hyperparameters such as `kernel`, and `random_state` to `linear`, and 0 respectively. The remaining hyperparameters of the support vector machine algorithm are set to default values.

```

"""

from sklearn.naive_bayes import GaussianNB

classifier = GaussianNB()

classifier.fit(X_train, y_train)

#Naive Bayes classifier model

GaussianNB(priors=None, var_smoothing=1e-09)

#Display the results (confusion matrix and accuracy)
"""

Here evaluation metrics such as confusion matrix and accuracy are used to evaluate the performance of
the model built using a decision tree classifier.

"""

from sklearn.metrics import confusion_matrix, accuracy_score

y_pred = classifier.predict(X_test)

cm = confusion_matrix(y_test, y_pred)

print(cm)

accuracy_score(y_test, y_pred)

```

```

=====
[[ 99   8]
 [  2  62]]
>

```

20.Future Sales Prediction using Python

```

"""

Created on Tue Jan 24 13:52:17 2023

```

@author: GEOMOL GEORGE

Future Sales Prediction using Python

```

"""

```

```

import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

import plotly.io as io

io.renderers.default='browser'

data = pd.read_csv("D:/GEO/BE COURSES/LAB/DATASET/futuresale prediction.csv")

```

```

print(data.head())

print(data.sample(5))
print(data.isnull().sum())
import plotly.express as px
import plotly.graph_objects as go
figure = px.scatter(data_frame = data, x="Sales",
                    y="TV", size="TV", trendline="ols")
figure.show()
figure = px.scatter(data_frame = data, x="Sales",
                    y="Newspaper", size="Newspaper", trendline="ols")
figure.show()
figure = px.scatter(data_frame = data, x="Sales",
                    y="Radio", size="Radio", trendline="ols")
figure.show()
correlation = data.corr()
print(correlation["Sales"].sort_values(ascending=False))
x = np.array(data.drop(["Sales"], 1))
y = np.array(data["Sales"])
xtrain, xtest, ytrain, ytest = train_test_split(x, y,
                                                test_size=0.2,
                                                random_state=42)

model = LinearRegression()
model.fit(xtrain, ytrain)
print(model.score(xtest, ytest))
features = [[TV, Radio, Newspaper]]
features = np.array([[230.1, 37.8, 69.2]])
print(model.predict(features))

```

```

===== I
      TV  Radio  Newspaper  Sales
0  230.1  37.8      69.2    22.1
1   44.5  39.3      45.1    10.4
2   17.2  45.9      69.3    12.0
3  151.5  41.3      58.5    16.5
4  180.8  10.8      58.4    17.9
      TV  Radio  Newspaper  Sales
5     8.7  48.9      75.0     7.2
101  296.4  36.3     100.9    23.8
110  225.8   8.2      56.5    18.4
154  187.8  21.1       9.5    20.6
53  182.6  46.2      58.7    21.2
TV      0
Radio    0
Newspaper  0
Sales    0
dtype: int64

```